

# **MACHINE LEARNING PROJECT**

## **PROJECT TITLE :**

**NEURAL STYLE TRANSFER**

## **TEAM MEMBERS :**

| <b>S No</b> | <b>Roll No</b> | <b>Name</b>       |
|-------------|----------------|-------------------|
| 1           | 2019103518     | M Fayeka Masoodha |
| 2           | 2019103584     | Pooja S           |
| 3           | 2019103601     | Gunnala Hema      |
| 4           | 2019103530     | Jeyasri Meenakshi |

## **1.TITLE OF THE PROJECT : NEURAL STYLE TRANSFER**

### **2.PROBLEM STATEMENT :**

Neural Style Transfer is a technique that allows us to generate an image with the same "content" as a base image, but with the "style" of our chosen picture.

Style transfer is extensively used in photo and video editing software. These deep learning approaches and professional style transfer models can easily be applied to devices like mobile phones and give users a real-time ability to style images and videos. Style transfer also provides new techniques that can change the way we look and deal with art. It makes high-rated and over-priced artistic work reproducible for office and home decor, or for advertisements. Transfer models may help us commercialize art. There are also many cloud-powered video game streams that use image style transfer.

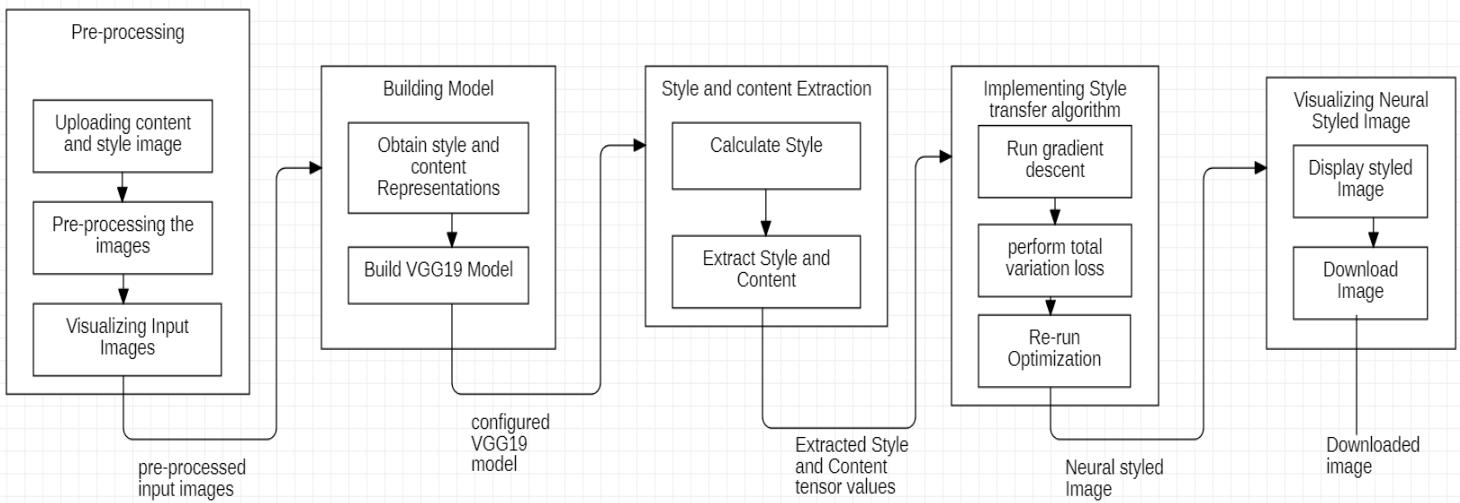
### **3.OBJECTIVE:**

Neural style transfer (NST) refers to a class of software algorithms that manipulate digital images, or videos, in order to adopt the appearance or visual style of another image. NST algorithms are characterized by their use of deep neural networks for the sake of image transformation. Common uses for NST are the creation of artificial artwork from photographs, for example by transferring the appearance of famous paintings to user-supplied photographs. Several notable mobile apps use NST techniques for this purpose, including DeepArt and Prisma. This method has been used by artists and designers around the globe to develop new artwork based on existent style.

Given a training pair of images—a photo and an artwork depicting that photo—a transformation could be learned and then applied to create new artwork from a new photo, by analogy.

## 4.OVERALL ARCHITECTURE:

NEURAL STYLE TRANSFER BLOCK DIAGRAM



## 5.MODULES

### I. PRE-PROCESSING

#### a. Uploading content and style image

- Importing and configuring the required modules like tensorflow, numpy, matplotlib etc.
- Obtain the content Image and style image.

### **b. Pre-process the images**

- Pre-processing involves detecting the image extension and converting it into a Tensor of required type.

### **c. Visualizing input images**

- Displaying the input images.

**Input :** Style and content Images

**Output :** Pre-processed Style and content Images.

## **II. BUILDING MODEL**

### **a. Obtain style and content representations**

- Loading the pre-trained arbitrary image stylization model from tensorflow hub.
- Intermediate layers of the model are used to get the content and style representations of the image.
- Here, the VGG19 network architecture which is a pretrained image classification network is used.

### **b. Build the VGG19 model**

- Load the VGG19 model and test run it on our image to ensure it's used correctly.
- Choose the intermediate layers from the network to represent the style and content of the image and provide it as an input to the VGG19 model.
- The VGG19 model then returns a list of intermediate layer outputs.

**Input :** Pre-processed Style and content Images

**Output :** configured VGG19 model

### **III. STYLE AND CONTENT EXTRACTION**

#### **a. Calculate Style**

- The content of an image is represented by the values of the intermediate feature maps.
- The style of an image can be described by the means and correlations across the different feature maps.
- Calculate the Gram matrix for a particular layer using the given formula.

$$G_{cd}^l = \frac{\sum_{ij} F_{ijc}^l(x) F_{ijd}^l(x)}{IJ}$$

#### **b. Extract style and content**

- Build a model that returns the style and content tensors.
- When called on an image, this model returns the gram matrix (style) of the style\_layers and content of the content\_layers.

**Input :** Configured VGG19 model

**Output :** Extracted style and content tensors.

## **IV. IMPLEMENT STYLE TRANSFER ALGORITHM**

### **a. Run Gradient Descent**

- With the above mentioned style and content extractor, implementing the style transfer algorithm.
- Calculating the mean square error for the image's output relative to each target, then taking the weighted sum of these losses.
- Create an optimizer. Optimization is then performed using the weighted combination of the two losses . The image is then updated and tested.

### **b. Perform total variation loss**

- One downside to this basic implementation is that it produces a lot of high frequency artifacts. We can decrease these by using an explicit regularization term on the high frequency components of the image.This is total variation loss.

### **c. Re-run the optimization**

- Choosing a weight for total variation loss
- Including it in the train step function.
- Reinitializing the optimization variable and running the optimization.

**Input :** style and content tensors

**Output :** neural styled image

## **V. VISUALIZING NEURAL STYLED IMAGE**

### **a. Display styled image:**

- The generated neural style clear output image can be displayed to the user.

### **b. Download image:**

- The final image will be saved in .png file format and given a name.
- Can be downloaded in the user's system.

**Input :** neural styled image

**Output :** Downloaded neural styled image

## **6. IMPLEMENTATION WITH SNAPSHOTS:**

### **I. PRE-PROCESSING**

#### **a. Uploading content and style image**



The screenshot shows a Jupyter Notebook interface with the following code in cell [1]:

```
[1] import os
import tensorflow as tf
# Load compressed models from tensorflow_hub
#The tensorflow_hub library currently supports two modes for downloading models.
#By default, a model is downloaded as a compressed archive and cached on disk.
os.environ['TFHUB_MODEL_LOAD_FORMAT'] = 'COMPRESSED'
```

The screenshot shows a Jupyter Notebook interface with the following code in cell [2]:

```
[2] import IPython.display as display

import matplotlib.pyplot as plt
import matplotlib as mpl
#setting up the size of the matplotlib image outputs
mpl.rcParams['figure.figsize'] = (12, 12)
#Switching off grid lines
mpl.rcParams['axes.grid'] = False

import numpy as np
import PIL.Image
```

At the bottom of the interface, it says "3s completed at 13:37".

```

+ Code + Text
✓ [2] import PIL.Image
✓ [3] import time
✓ [3] import functools
{x} ✓ [3]
[3] def tensor_to_image(tensor):
    tensor = tensor*255
    tensor = np.array(tensor, dtype=np.uint8)
    if np.ndim(tensor)>3:
        assert tensor.shape[0] == 1
        tensor = tensor[0]
    return PIL.Image.fromarray(tensor)

[4] content_path = tf.keras.utils.get_file('YellowLabradorLooking_new.jpg', 'https://storage.googleapis.com/download.tensorflow.org/example_images/YellowLabradorLooking_new.jpg')
style_path = tf.keras.utils.get_file('kandinsky5.jpg','https://storage.googleapis.com/download.tensorflow.org/example_images/Vassily_Kandinsky%2C_1913.jpg')

Downloading data from https://storage.googleapis.com/download.tensorflow.org/example_images/YellowLabradorLooking_new.jpg
90112/83281 [=====] - 0s 0us/step
98304/83281 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/download.tensorflow.org/example_images/Vassily_Kandinsky%2C_1913 - Composition_7.jpg
196608/195196 [=====] - 0s 0us/step
204800/195196 [=====] - 0s 0us/step

```

## b. Pre-process the images

**PRE-PROCESSING INPUT IMAGES**

```

+ Code + Text
{x} ✓ [0] def load_img(path_to_img):
    max_dim = 512
    #reads file
    img = tf.io.read_file(path_to_img)
    print(img)
    img = tf.image.decode_image(img, channels=3)
    img = tf.image.convert_image_dtype(img, tf.float32)

    shape = tf.cast(tf.shape(img)[-1], tf.float32)
    long_dim = max(shape)
    scale = max_dim / long_dim

    new_shape = tf.cast(shape * scale, tf.int32)

    img = tf.image.resize(img, new_shape)
    img = img[tf.newaxis, :]
    return img

[6] def imshow(image, title=None):
    if len(image.shape) > 3:

```

✓ 3s completed at 13:37

```

+ Code + Text
{x} ✓ [0] #reads file
[5] [5] img = tf.io.read_file(path_to_img)
print(img)
img = tf.image.decode_image(img, channels=3)
img = tf.image.convert_image_dtype(img, tf.float32)

shape = tf.cast(tf.shape(img)[-1], tf.float32)
long_dim = max(shape)
scale = max_dim / long_dim

new_shape = tf.cast(shape * scale, tf.int32)

img = tf.image.resize(img, new_shape)
img = img[tf.newaxis, :]
return img

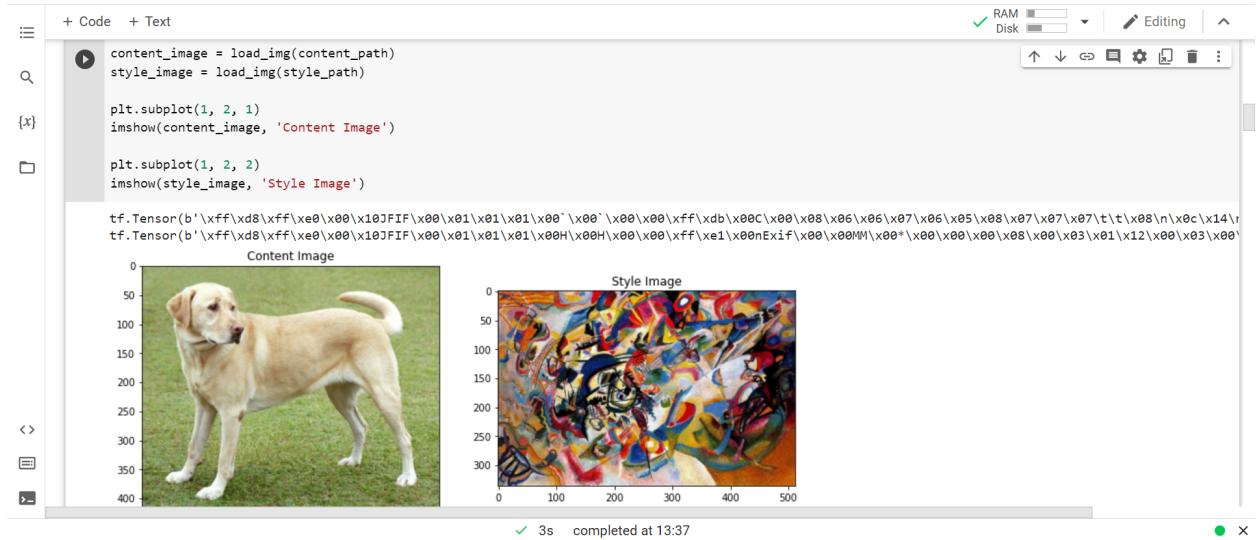
[6] def imshow(image, title=None):
    if len(image.shape) > 3:
        image = tf.squeeze(image, axis=0)

    plt.imshow(image)
    if title:
        plt.title(title)

```

✓ 3s completed at 13:37

## c. Visualizing input images

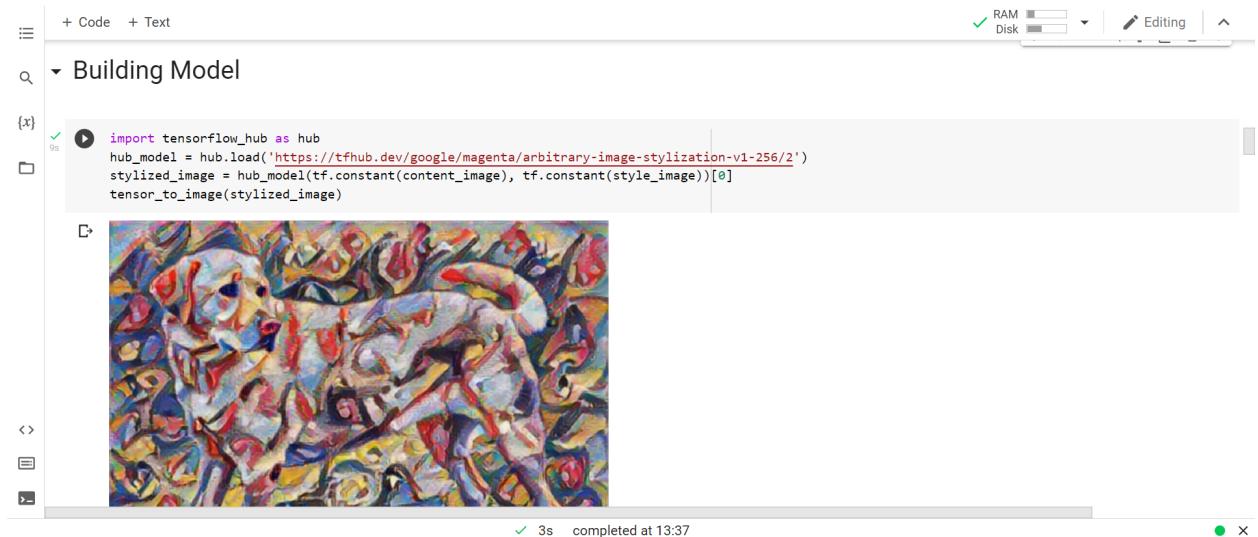


```
+ Code + Text  
content_image = load_img(content_path)  
style_image = load_img(style_path)  
  
plt.subplot(1, 2, 1)  
imshow(content_image, 'Content Image')  
  
plt.subplot(1, 2, 2)  
imshow(style_image, 'Style Image')  
  
Content Image  
Style Image
```

3s completed at 13:37

## II. BUILDING MODEL

### a. Obtain style and content representations



```
+ Code + Text  
Building Model  
import tensorflow_hub as hub  
hub_model = hub.load('https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/2')  
stylized_image = hub_model(tf.constant(content_image), tf.constant(style_image))[0]  
tensor_to_image(stylized_image)
```

3s completed at 13:37

+ Code + Text

✓ [9] x = tf.keras.applications.vgg19.preprocess\_input(content\_image\*255)  
x = tf.image.resize(x, (224, 224))  
vgg = tf.keras.applications.VGG19(include\_top=True, weights='imagenet')  
{x}  
prediction\_probabilities = vgg(x)  
prediction\_probabilities.shape

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels.h5)  
574717952/574710816 [=====] - 8s 0us/step  
574726144/574710816 [=====] - 8s 0us/step  
TensorShape([1, 1000])

✓ [10] predicted\_top\_5 = tf.keras.applications.vgg19.decode\_predictions(prediction\_probabilities.numpy())[0]  
[(class\_name, prob) for (number, class\_name, prob) in predicted\_top\_5]

Downloading data from [https://storage.googleapis.com/download.tensorflow.org/data/imagenet\\_class\\_index.json](https://storage.googleapis.com/download.tensorflow.org/data/imagenet_class_index.json)  
40960/35363 [=====] - 0s 0us/step  
49152/35363 [=====] - 0s 0us/step  
[('Labrador\_retriever', 0.4931729),  
 ('golden\_retriever', 0.23665176),  
 ('kuvasz', 0.036357265),  
 ('Chesapeake\_Bay\_retriever', 0.024182765),  
 ('Greater\_Swiss\_Mountain\_dog', 0.018646065)]

✓ [11] vgg = tf.keras.applications.VGG19(include\_top=False, weights='imagenet')

✓ 3s completed at 13:37

+ Code + Text

✓ [11] vgg = tf.keras.applications.VGG19(include\_top=False, weights='imagenet')  
{x}  
print()  
for layer in vgg.layers:  
 print(layer.name)

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5)  
80142336/80134624 [=====] - 1s 0us/step  
80150528/80134624 [=====] - 1s 0us/step

input\_2  
block1\_conv1  
block1\_conv2  
block1\_pool  
block2\_conv1  
block2\_conv2  
block2\_pool  
block3\_conv1  
block3\_conv2  
block3\_conv3  
block3\_conv4  
block3\_pool  
block4\_conv1  
block4\_conv2  
block4\_conv3  
block4\_conv4

✓ 3s completed at 13:37

```
+ Code + Text
[11] block3_conv3
block3_conv4
block3_pool
block4_conv1
block4_conv2
block4_conv3
block4_conv4
block4_pool
block5_conv1
block5_conv2
block5_conv3
block5_conv4
block5_pool

[12] content_layers = ['block5_conv2']

style_layers = ['block1_conv1',
                'block2_conv1',
                'block3_conv1',
                'block4_conv1',
                'block5_conv1']

num_content_layers = len(content_layers)
num_style_layers = len(style_layers)

✓ 3s completed at 13:37
```

## b. Build the VGG19 model

```
+ Code + Text
Build model

[13] def vgg_layers(layer_names):
    """ Creates a vgg model that returns a list of intermediate output values."""
    # Load our model. Load pretrained VGG, trained on imagenet data
    vgg = tf.keras.applications.VGG19(include_top=False, weights='imagenet')
    vgg.trainable = False

    outputs = [vgg.get_layer(name).output for name in layer_names]

    model = tf.keras.Model([vgg.input], outputs)
    return model

#Create model
style_extractor = vgg_layers(style_layers)
style_outputs = style_extractor(style_image*255)

#Look at the statistics of each layer's output
for name, output in zip(style_layers, style_outputs):
    print(name)
    print(" shape: ", output.numpy().shape)
    print(" min: ", output.numpy().min())
    print(" max: ", output.numpy().max())
    print(" mean: ", output.numpy().mean())
    print()

✓ 3s completed at 13:37
```

```
+ Code + Text
[14] print(name)
print(" shape: ", output.numpy().shape)
print(" min: ", output.numpy().min())
print(" max: ", output.numpy().max())
print(" mean: ", output.numpy().mean())
print()

block1_conv1
shape: (1, 336, 512, 64)
min: 0.0
max: 835.5255
mean: 33.97525

block2_conv1
shape: (1, 168, 256, 128)
min: 0.0
max: 4625.887
mean: 199.82687

block3_conv1
shape: (1, 84, 128, 256)
min: 0.0
max: 8789.24
mean: 230.78099

block4_conv1
shape: (1, 42, 64, 512)
```

```

+ Code + Text
[14]   mean: 33.97525
      block2_conv1
      shape: (1, 168, 256, 128)
      min: 0.0
      max: 4625.887
      mean: 199.82687
      block3_conv1
      shape: (1, 84, 128, 256)
      min: 0.0
      max: 8789.24
      mean: 230.78099
      block4_conv1
      shape: (1, 42, 64, 512)
      min: 0.0
      max: 21566.135
      mean: 791.24005
      block5_conv1
      shape: (1, 21, 32, 512)
      min: 0.0
      max: 3189.2537
      mean: 59.179478

```

✓ 3s completed at 13:37

### III. STYLE AND CONTENT EXTRACTION

#### a. Calculate Style

```

+ Code + Text
Calculate style
[15] def gram_matrix(input_tensor):
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensor, input_tensor)
    input_shape = tf.shape(input_tensor)
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)
    return result/(num_locations)

```

#### b. Extract style and content

```

+ Code + Text
Extract style and content
[16] class StyleContentModel(tf.keras.models.Model):
    def __init__(self, style_layers, content_layers):
        super(StyleContentModel, self).__init__()
        self.vgg = vgg19.VGG19(style_layers + content_layers)
        self.style_layers = style_layers
        self.content_layers = content_layers
        self.num_style_layers = len(style_layers)
        self.vgg.trainable = False

    def call(self, inputs):
        "Expects float input in [0,1]"
        inputs = inputs*255.0
        preprocessed_input = tf.keras.applications.vgg19.preprocess_input(inputs)
        outputs = self.vgg(preprocessed_input)
        style_outputs, content_outputs = (outputs[:self.num_style_layers],
                                         outputs[self.num_style_layers:])

    style_outputs = [gram_matrix(style_output)
                    for style_output in style_outputs]

```

✓ 3s completed at 13:37

```
+ Code + Text
[16]     content_dict = {content_name: value
                     for content_name, value
                     in zip(self.content_layers, content_outputs)}
{x}
style_dict = {style_name: value
              for style_name, value
              in zip(self.style_layers, style_outputs)}

return {'content': content_dict, 'style': style_dict}

[17] extractor = StyleContentModel(style_layers, content_layers)

results = extractor(tf.constant(content_image))

print('Styles')
for name, output in sorted(results['style'].items()):
    print(" ", name)
    print(" shape: ", output.numpy().shape)
    print(" min: ", output.numpy().min())
    print(" max: ", output.numpy().max())
    print(" mean: ", output.numpy().mean())
    print()

print("Contents:")

```

✓ 3s completed at 13:37

```
+ Code + Text
[17] print("Contents:")
for name, output in sorted(results['content'].items()):
    print(" ", name)
    print(" shape: ", output.numpy().shape)
    print(" min: ", output.numpy().min())
    print(" max: ", output.numpy().max())
    print(" mean: ", output.numpy().mean())

    Styles:
        block1_conv1
        shape: (1, 64, 64)
        min: 0.0055228495
        max: 28014.562
        mean: 263.79022

        block2_conv1
        shape: (1, 128, 128)
        min: 0.0
        max: 61479.484
        mean: 9100.949

        block3_conv1
        shape: (1, 256, 256)
        min: 0.0
        max: 545623.4
        mean: 7660.976

```

✓ 3s completed at 13:37

```
+ Code + Text
[17]
mean: 9100.949
block3_conv1
shape: (1, 256, 256)
min: 0.0
max: 545623.4
mean: 7660.976

block4_conv1
shape: (1, 512, 512)
min: 0.0
max: 4320501.0
mean: 134288.86

block5_conv1
shape: (1, 512, 512)
min: 0.0
max: 110005.37
mean: 1487.0378

Contents:
block5_conv2
shape: (1, 26, 32, 512)
min: 0.0
max: 2410.8794
mean: 13.764149

```

✓ 3s completed at 13:37

## IV. IMPLEMENT STYLE TRANSFER ALGORITHM

### a. Run Gradient Descent

```
+ Code + Text RAM Disk Editing ^  
Run gradient descent  
  
{x} [18] style_targets = extractor(style_image)['style']  
content_targets = extractor(content_image)['content']  
[x] [19] image = tf.Variable(content_image)  
[x] [20] def clip_0_1(image):  
    return tf.clip_by_value(image, clip_value_min=0.0, clip_value_max=1.0)  
[x] [21] opt = tf.optimizers.Adam(learning_rate=0.02, beta_1=0.99, epsilon=1e-1)  
[x] [22] style_weight=1e-2  
      content_weight=1e4  
[x] [23] def style_content_loss(outputs):  
    style_outputs = outputs['style']  
    content_outputs = outputs['content']  
    style_loss = tf.add_n([tf.reduce_mean((style_outputs[name]-style_targets[name])**2)  
    for name in style_outputs.keys()])  
    style_loss *= style_weight / num_style_layers  
    content_loss = tf.add_n([tf.reduce_mean((content_outputs[name]-content_targets[name])**2)  
    for name in content_outputs.keys()])  
    content_loss *= content_weight / num_content_layers  
    loss = style_loss + content_loss  
    return loss  
[x] [24] @tf.function()  
def train_step(image):  
    with tf.GradientTape() as tape:  
        outputs = extractor(image)  
        loss = style_content_loss(outputs)  
        grad = tape.gradient(loss, image)  
        opt.apply_gradients([(grad, image)])  
        image.assign(clip_0_1(image))  
[x] [25] train_step(image)  
train step(image)  
✓ 3s completed at 13:37
```

+ Code + Text

✓ [25] train\_step(image)  
train\_step(image)  
tensor\_to\_image(image)

{x}

🕒 RAM Disk Editing ^



3s completed at 13:37

+ Code + Text

✓ [26] import time  
start = time.time()

{x} epochs = 10  
steps\_per\_epoch = 100

🕒 RAM Disk Editing ^

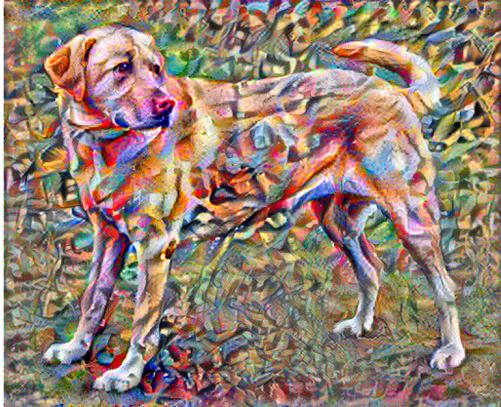
```
for n in range(epochs):
    for m in range(steps_per_epoch):
        step += 1
        train_step(image)
        print(".", end='', flush=True)
        display.clear_output(wait=True)
        display.display(tensor_to_image(image))
        print("Train step: {}".format(step))

end = time.time()
print("Total time: {:.1f}".format(end-start))
```

+ Code + Text

✓ [26]

🕒 RAM Disk Editing ^



Train step: 500

3s completed at 13:37

## b. Perform total variation loss

+ Code + Text

### Total variation loss

```
{x} [27] def high_pass_x_y(image):
    x_var = image[:, :, 1:, :] - image[:, :, :-1, :]
    y_var = image[:, 1:, :, :] - image[:, :-1, :, :]

    return x_var, y_var

[28] x_deltas, y_deltas = high_pass_x_y(content_image)

plt.figure(figsize=(14, 10))
plt.subplot(2, 2, 1)
imshow(clip_0_1(2*y_deltas+0.5), "Horizontal Deltas: Original")

plt.subplot(2, 2, 2)
imshow(clip_0_1(2*x_deltas+0.5), "Vertical Deltas: Original")

<> x_deltas, y_deltas = high_pass_x_y(image)

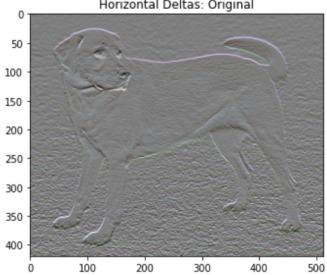
plt.subplot(2, 2, 3)
imshow(clip_0_1(2*y_deltas+0.5), "Horizontal Deltas: Styled")
```

✓ 3s completed at 13:37

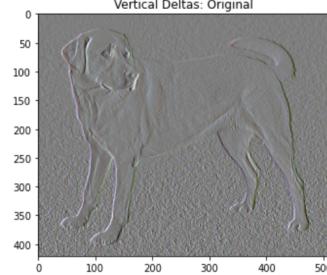
+ Code + Text

```
[28] plt.subplot(2, 2, 4)
imshow(clip_0_1(2*x_deltas+0.5), "Vertical Deltas: Styled")
```

{x}



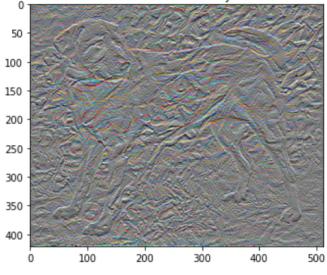
Horizontal Deltas: Original



Vertical Deltas: Original

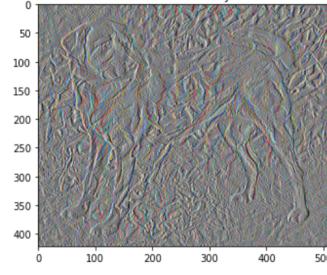
+ Code + Text

```
[28] Horizontal Deltas: Styled
```



Horizontal Deltas: Styled

```
[28] Vertical Deltas: Styled
```



Vertical Deltas: Styled

✓ RAM Disk Editing ^

```
[ ] def total_variation_loss(image):
    x_deltas, y_deltas = high_pass_x_y(image)
    return tf.reduce_sum(tf.abs(x_deltas)) + tf.reduce_sum(tf.abs(y_deltas))

[ ] total_variation_loss(image).numpy()

149353.12

[ ] tf.image.total_variation(image).numpy()

array([149353.12], dtype=float32)
```

## c. Re-run the optimization

Re-run the optimization

```
[ ] total_variation_weight=30

❶ @tf.function()
def train_step(image):
    with tf.GradientTape() as tape:
        outputs = extractor(image)
        loss = style_content_loss(outputs)
        loss += total_variation_weight*tf.image.total_variation(image)

    grad = tape.gradient(loss, image)
    opt.apply_gradients([(grad, image)])
    image.assign(clip_0_1(image))

[ ] image = tf.Variable(content_image)
```

```
✓ ❷ import time
dm start = time.time()

epochs = 10
steps_per_epoch = 100

step = 0
for n in range(epochs):
    for m in range(steps_per_epoch):
        step += 1
        train_step(image)
        print(".", end='', flush=True)
        display.clear_output(wait=True)
        display.display(tensor_to_image(image))
        print("Train step: {}".format(step))

end = time.time()
print("Total time: {:.1f}".format(end-start))
```

## **V. VISUALIZING NEURAL STYLED IMAGE**

### **a. Display styled image**



Train step: 1000  
Total time: 259.1

### **b. Download image**

```
file_name = 'stylized_image.png'
tensor_to_image(image).save(file_name)

try:
    from google.colab import files
except ImportError:
    pass
else:
    files.download(file_name)
```

stylized\_image.png

## 7.DATASETS

- In this project VGG19 model, with weights pre-trained on ImageNet is used.
- ImageNet, is a dataset of over 15 millions labeled high-resolution images with around 22,000 categories.
- Also as input a style image and a content image is uploaded.

## 8.PERFORMANCE MEASURES

### I.Content Loss

Let p and x be the original image and the image that is generated and P and F their respective feature representation in layer i .The squared-error loss between the two feature representations is defined as the content loss.

$$\mathcal{L}_{content}(\vec{p}, \vec{x}, l) = \frac{1}{2} \sum_{i,j} (F_{ij}^l - P_{ij}^l)^2$$

### II.Style Loss

let a and x be the original image and the image that is generated and Ai and Gi their respective style representations in layer i . The contribution of that layer to the total loss is then

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{ij}^l - A_{ij}^l)^2$$

$$\mathcal{L}_{style}(\vec{a}, \vec{x}) = \sum_{l=0}^L w_l E_l$$

G - Gram matrix

$$G_{ij}^l = \sum_k F_{ik}^l F_{jk}^l.$$

### III.Cost Function For Style Transfer

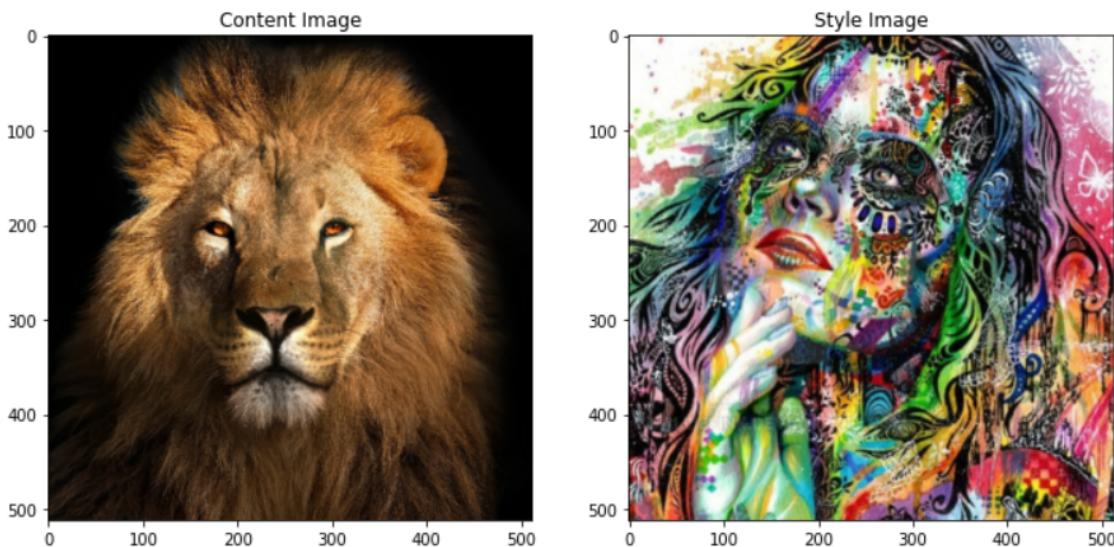
The total cost function is the weighted sum of two separate costs namely, style cost and content cost calculated separately.

$$J_{total}(G) = \alpha \times J_{content}(C, G) + \beta \times J_{style}(S, G)$$

Here,  $\alpha$  is called the content weight and  $\beta$  is called the style weight. Generally, style weight is much higher than the content weight as to emphasize more on the “style transfer”.

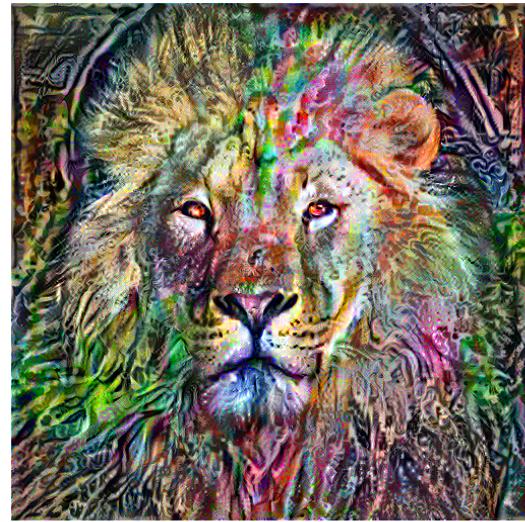
## 9.TEST CASE:

### a) Content image and Style image

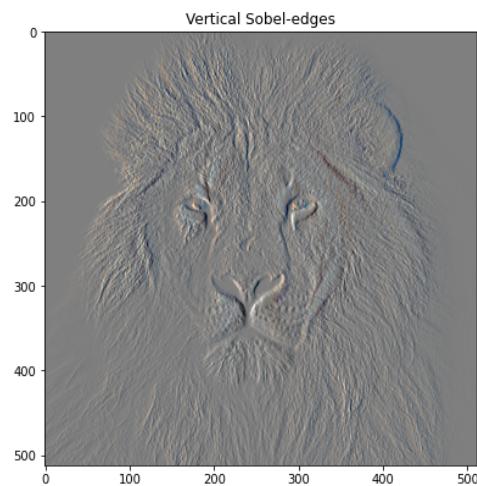
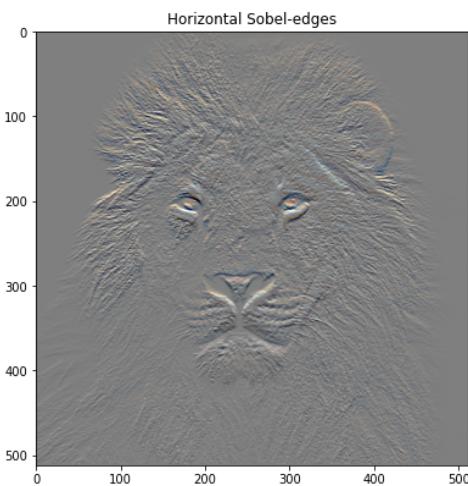
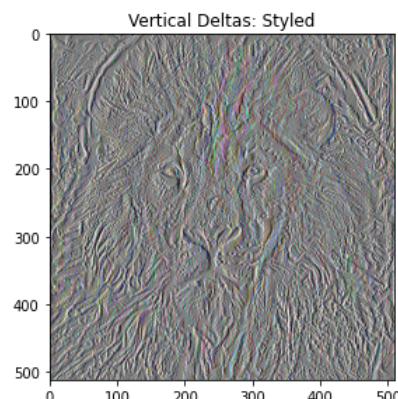
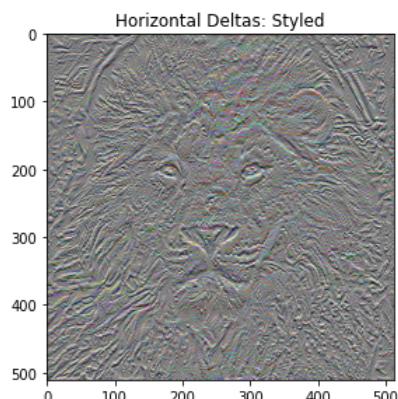
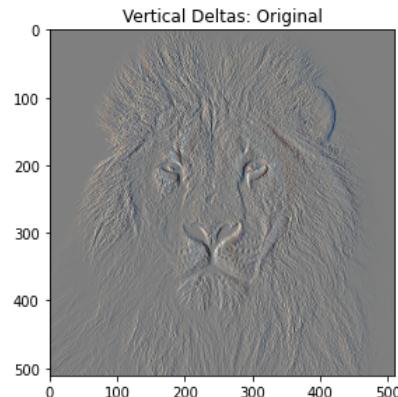
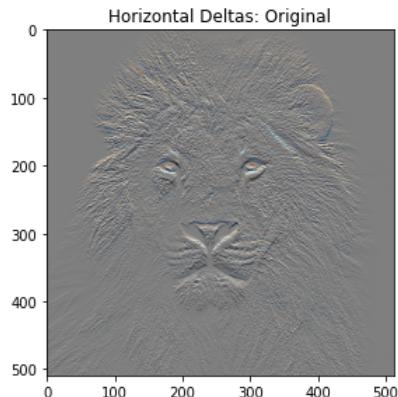


### Implementing style transfer algorithm – Outputs

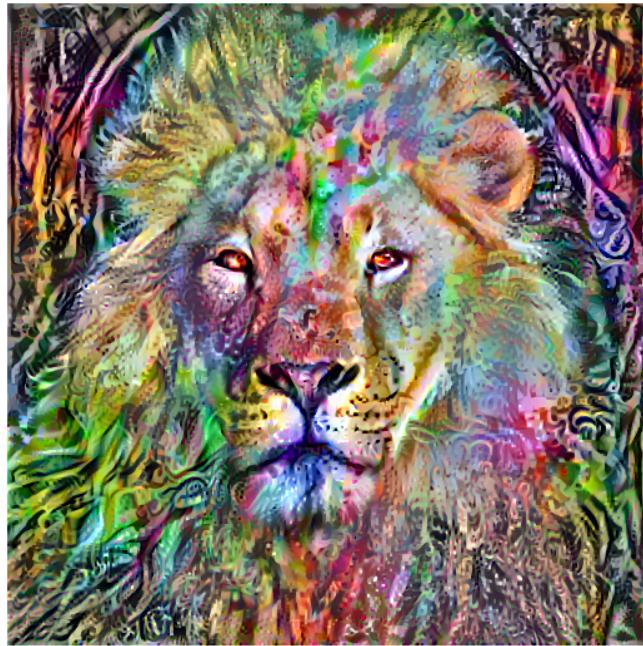
#### On Running Gradient Descent



## On performing total variation loss function:



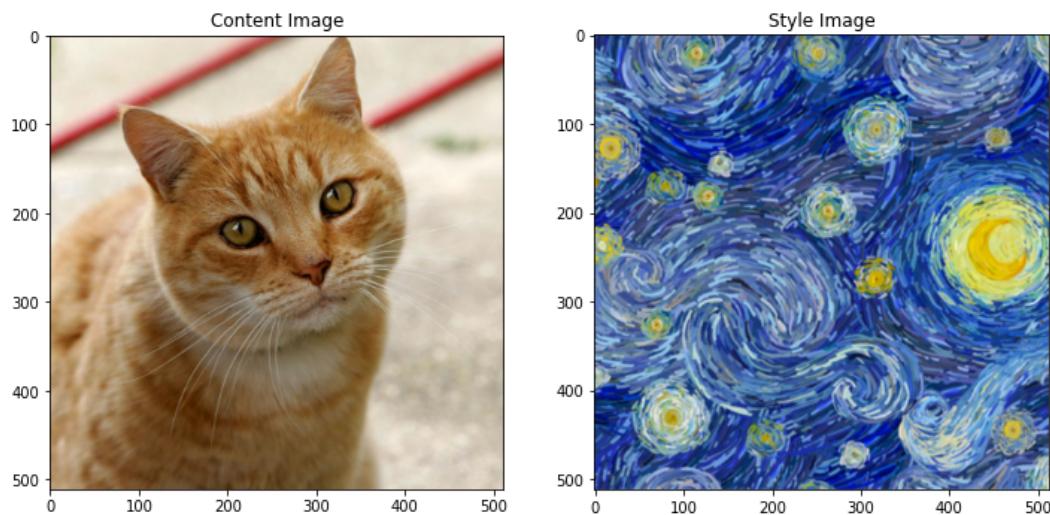
**After re-running the optimization  
Final stylized image :-**



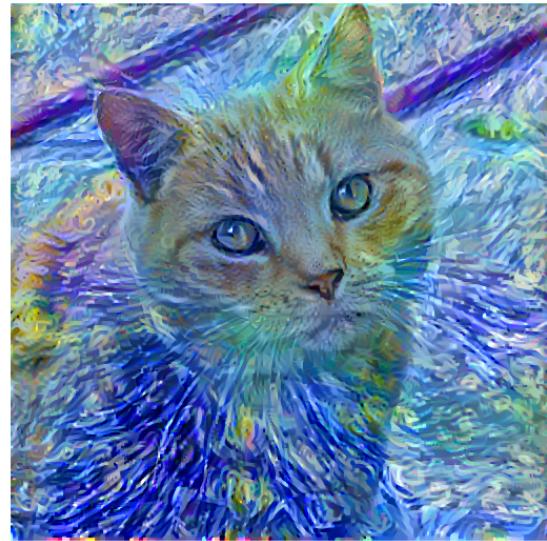
Train step: 1000  
Total time: 8281.6

**b) Content image and Style image**

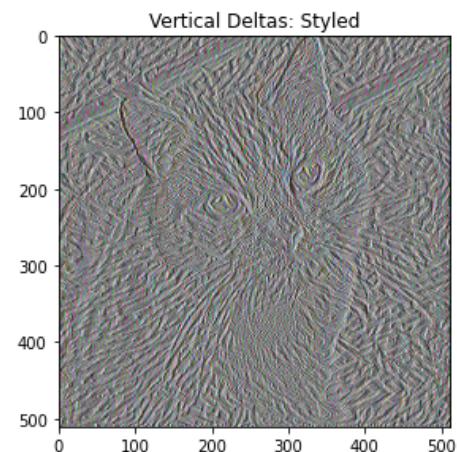
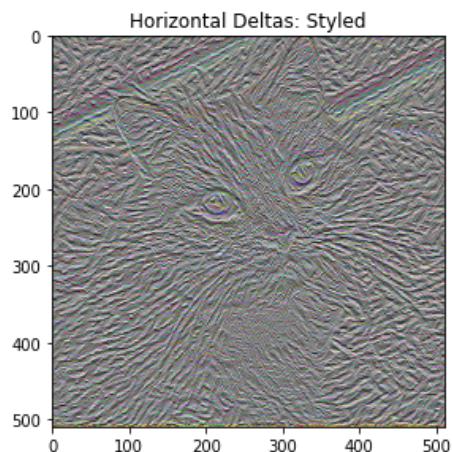
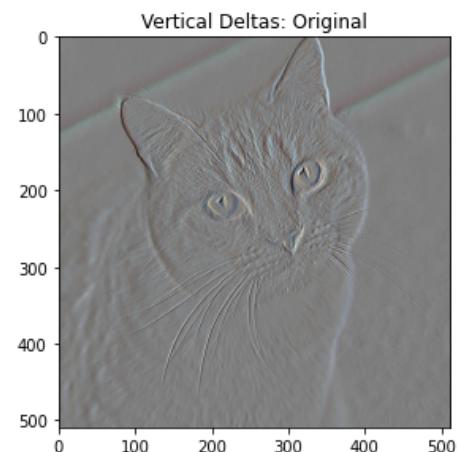
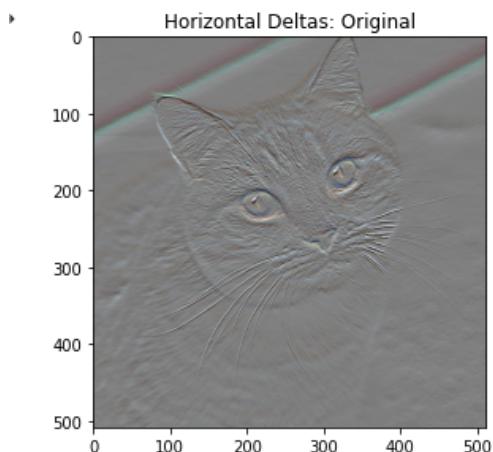
**Implementing style transfer algorithm – Outputs**

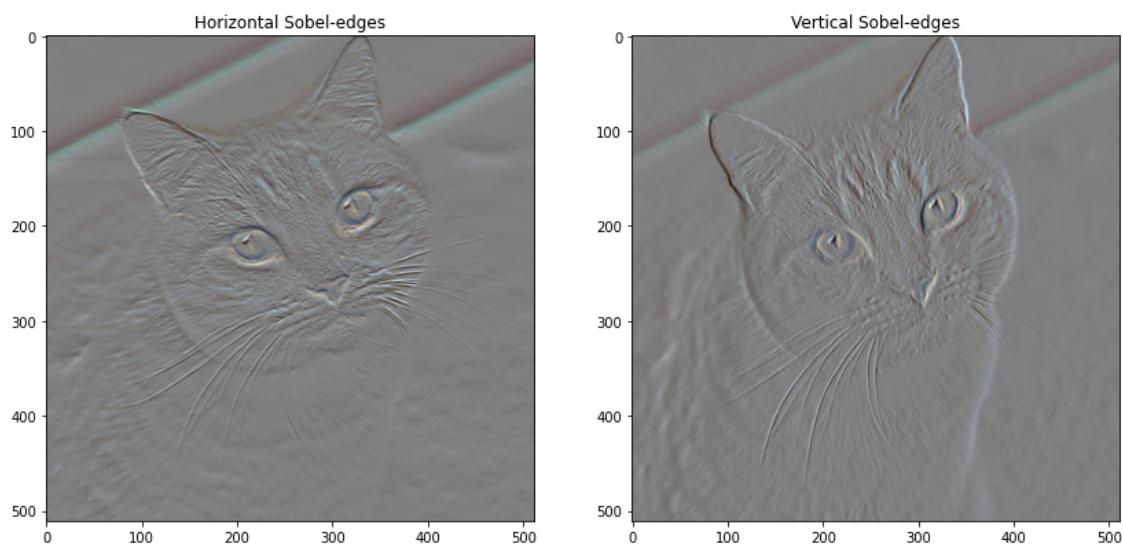


## On Running Gradient Descent

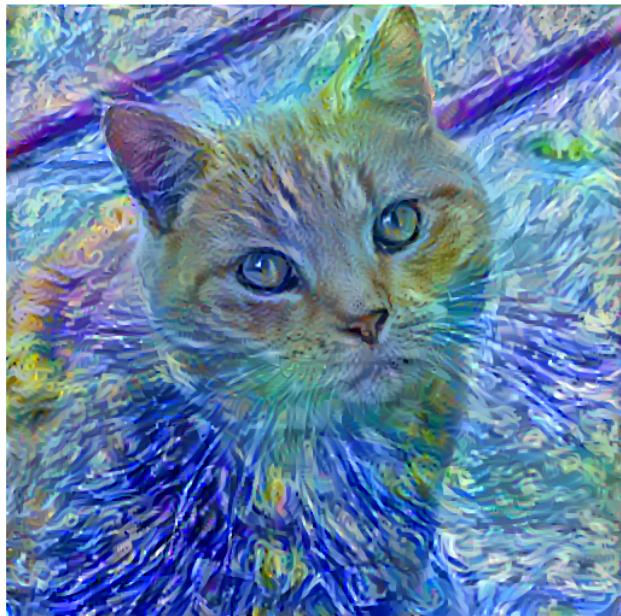


## On performing total variation loss function:





## After re-running the optimization Final stylized image :-



Train step: 1000  
Total time: 6113.3

## **10.EVALUATION**

The most common approach in quantitative evaluation is to compute the runtime or the convergence speed(how long it takes the loss function to converge) of different algorithms.

### **For the test cases:**

1) For 1000 training steps

Run time : 6113.3 seconds (101 minutes

approximately)

2) For 1000 training steps

Run time : 8281.6 seconds (138 minutes  
approximately)

## **11.REFERENCES**

- [Neural Style Transfer: Everything You Need to Know \[Guide\] \(v7labs.com\)](#)
- <https://sci-hub.hkvisa.net/https://ieeexplore.ieee.org/document/9275956>
- [Style in Computer Vision — Neural Style Transfer | by Siddhartha Gairola | DataDrivenInvestor](#)